

Direct Head-to-Head: RLSTCcode vs Q-RLSTC

Metric Definition

Competitive Ratio (CR) measures segmentation-induced clustering quality relative to an unsegmented baseline:

$$CR = \frac{OD_{\text{segmented}}}{OD_{\text{baseline}}}$$

where OD is the average over-distance (mean IED from each sub-trajectory to its assigned cluster centre). Lower CR is better: CR < 1.0 means the RL agent produces tighter clusters than the baseline; CR = 1.0 is parity; CR > 1.0 is worse.

All Models Ranked by Competitive Ratio

Rank	Model	System	Params	Val CR	Optimizer	Training Data
1	RLSTCcode (best, 3k)	RLSTCcode	514	0.5892	SGD	3,000 trajs
2	RLSTCcode (5-fold CV)	RLSTCcode	514	0.7543 +/- 0.0369	SGD	~4,000 trajs
3	VQ-DQN (5q x 3L)	Q-RLSTC	34	1.4811	SPSA	30 trajs
4	RLSTCcode (model-state4)	RLSTCcode	514	1.5453	SGD	unknown
5	Control C (h=32x32)	Q-RLSTC	1,314	1.6390	SPSA	30 trajs
6	Control B (h=64)	Q-RLSTC	514	5.1133	SPSA	30 trajs
7	Control A (linear)	Q-RLSTC	12	9.2051	SPSA	30 trajs

Uncertainty note: RLSTCcode reports 5-fold CV variance. Q-RLSTC results are single-seed (seed=42). Multi-seed confidence intervals for the SPSA models are a required follow-up (see Limitations).

Claim A: Absolute Clustering Quality

RLSTCcode achieves the best absolute CR when given sufficient data and gradient access.

Dimension	RLSTCcode (best)	VQ-DQN
Val CR	0.5892	1.4811
Parameters	514	34
Training data	3,000 trajs	30 trajs
Optimizer	SGD (backprop)	SPSA (gradient-free)
CR ratio	1.0x	2.5x worse

Dimension	RLSTCcode (best)	VQ-DQN
Data ratio	1.0x	100x less

RLSTCcode wins on absolute CR (0.5892 vs 1.4811) thanks to 100x more training data and gradient-based optimization. This comparison conflates three variables (optimizer, data scale, model class) and is not a controlled experiment.

Claim B: Parameter-Efficient Inductive Bias under SPSA

Under identical gradient-free training conditions, the 34-parameter VQ-DQN outperforms all classical controls including one with 39x more parameters.

Controlled Conditions (identical for all four models)

Dimension	Value
Optimizer	SPSA (gradient-free, shared implementation)
Training data	30 trajectories (27 train / 3 val)
Epochs	2
Batch size	32
Replay buffer	5,000
Gamma	0.9
Epsilon schedule	1.0 -> 0.1 (decay 0.99)
Target update	Hard copy every 10 episodes
Double DQN	Yes
Seed	42

SPSA Hyperparameters (shared by all models)

Parameter	Value	Role
A (stability)	20	Warmup period for learning rate
a (initial LR)	0.12	Learning rate scale
c (perturbation)	0.08	Finite-difference step size
alpha	0.602	LR decay exponent
gamma	0.101	Perturbation decay exponent
Grad clip	10.0	Max gradient norm
Momentum (m-SPSA)	Disabled for E1	No gradient averaging
Tuning budget	None	Defaults from Spall (1998)

Results

Model	Params	Val CR	CUT%	Segments	Wall Time
VQ-DQN (5q x 3L)	34	1.4811	33.2%	465	987s
Control C (h=32x32)	1,314	1.6390	3.9%	57	10s
Control B (h=64)	514	5.1133	0.2%	6	9s
Control A (linear)	12	9.2051	0.0%	3	9s

VQ-DQN achieves 9.6% lower CR with 39x fewer parameters than the best classical control.

CUT% Interpretation

The VQ-DQN's 33.2% cut rate produces 465 segments from 30 trajectories (avg ~15.5 segments per trajectory). This is not degenerate over-segmentation: the D1 random-baseline experiment shows that random cutting at $p=0.5$ achieves $CR=1.4311$ with 453 segments (comparable segment count, slightly better CR). The VQ-DQN's learned policy achieves a similar segment count through selective, reward-driven cuts rather than uniform random ones.

For context, min-segment-length $L_{MIN}=3$ prevents trivially short segments, and the $CUT_PENALTY=0.12$ in the reward function explicitly penalises excessive cutting.

Same-Architecture Comparison (Dense 5->64->2, 514 params)

Model	Val CR	Optimizer	Data	Result
RLSTCcode	0.5892	SGD	3,000	Learns effective policy
Q-RLSTC Control B	5.1133	SPSA	30	Fails to learn cutting

The identical 514-param MLP degrades 8.7x when switching from SGD+3k data to SPSA+30 data. This confirms that the gap between Claim A and Claim B is driven by optimizer and data scale, not architecture.

Takeaway

With the same gradient-free optimizer and the same 30-trajectory training regime, a 34-parameter VQ-DQN matches or beats substantially larger classical controls, suggesting parameter-efficient inductive bias in the small-data SPSA setting — while classical SGD with sufficient data still wins on absolute CR.

Limitations

1. Multi-seed uncertainty: SPSA results are from a single seed (42). Confidence intervals across 5+ seeds are needed to confirm statistical significance of Claim B.
2. Data-scale confound: The 100x data gap between Claim A and Claim B prevents fair absolute comparison. Running RLSTCcode at 30 trajectories (or Q-RLSTC at 3,000) would isolate the model-class variable.
3. SPSA tuning: All models share default SPSA hyperparameters (Spall 1998). It is possible that architecture-specific tuning of (a, c, A) could improve classical control performance under SPSA.
4. Epoch count: 2 epochs is minimal. Extended training may narrow or widen the quantum-classical gap.

Plots




Figure 1: Unified Comparison




Figure 2: Data Scaling



Figure 3: Controlled Comparison




Figure 4: Parameter Efficiency




Figure 5: Summary Dashboard